

|   |
|---|
| Experiment No. 3  |
| Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model |
| Date of Performance: 16/08/2023   |
| Date of Submission: 13/09/2023  |

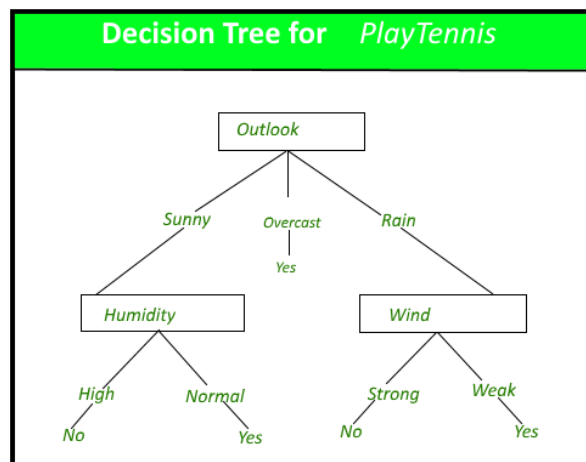


**Aim:** Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

**Objective:** To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

**Theory:**

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



**Dataset:**

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:



## Vidyavardhini's College of Engineering & Technology

### Department of Computer Engineering

---

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras,



Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

### **Conclusion:**

1. Discuss about the how categorical attributes have been dealt with during data pre-processing.

Categorical attributes, also known as categorical variables or features, are those that represent discrete and distinct categories or groups. These attributes are common in data across various domains, such as gender, color, city names, and more. Label encoding is a simple technique where each category is assigned a unique integer label. For example, if you have a "color" attribute with categories like "red," "blue," and "green," you can encode them as 0, 1, and 2, respectively. Here Label encoder is used to convert categorical to numerical.

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.

Hyperparameter tuning is a critical step in optimizing the performance of a decision tree model. Decision trees have several hyperparameters that control their structure and behavior, and tuning these hyperparameters can help improve the model's accuracy, generalization, and robustness. Max depth controls the maximum depth of the tree. A deeper tree can capture more complex relationships in the data but is prone to overfitting. We have used max depth as 5.

3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

Accuracy obtained in the decision tree model is 85%. A confusion matrix is a tabular representation used in machine learning to evaluate the performance of a classification model, especially for binary classification problems. Precision Obtain is 0.98 ,recall obtained is 0.27 and f1 score is 0.43.

## adultcensus

Adult Census Income Dataset

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
```

```
[ ]: df=pd.read_csv("/content/adult.csv",error_bad_lines=False, engine="python")
```

<ipython-input-265-8dbf926b9719>:1: FutureWarning: The error\_bad\_lines argument has been deprecated and will be removed in a future version. Use on\_bad\_lines in the future.

```
df=pd.read_csv("/content/adult.csv",error_bad_lines=False, engine="python")
```

```
[ ]: print(df)
```

|       | age | workclass | fnlwgt | education    | education.num | marital.status     | \ |
|-------|-----|-----------|--------|--------------|---------------|--------------------|---|
| 0     | 90  | ?         | 77053  | HS-grad      | 9             | Widowed            |   |
| 1     | 82  | Private   | 132870 | HS-grad      | 9             | Widowed            |   |
| 2     | 66  | ?         | 186061 | Some-college | 10            | Widowed            |   |
| 3     | 54  | Private   | 140359 | 7th-8th      | 4             | Divorced           |   |
| 4     | 41  | Private   | 264663 | Some-college | 10            | Separated          |   |
| ...   | ... | ...       | ...    | ...          | ...           | ...                |   |
| 32556 | 22  | Private   | 310152 | Some-college | 10            | Never-married      |   |
| 32557 | 27  | Private   | 257302 | Assoc-acdm   | 12            | Married-civ-spouse |   |
| 32558 | 40  | Private   | 154374 | HS-grad      | 9             | Married-civ-spouse |   |
| 32559 | 58  | Private   | 151910 | HS-grad      | 9             | Widowed            |   |
| 32560 | 22  | Private   | 201490 | HS-grad      | 9             | Never-married      |   |

|   | occupation        | relationship  | race  | sex    | capital.gain | \ |
|---|-------------------|---------------|-------|--------|--------------|---|
| 0 | ?                 | Not-in-family | White | Female | 0            |   |
| 1 | Exec-managerial   | Not-in-family | White | Female | 0            |   |
| 2 | ?                 | Unmarried     | Black | Female | 0            |   |
| 3 | Machine-op-inspct | Unmarried     | White | Female | 0            |   |

|       |                   |               |       |        |     |
|-------|-------------------|---------------|-------|--------|-----|
| 4     | Prof-specialty    | Own-child     | White | Female | 0   |
| ...   | ...               | ...           | ...   | ...    | ... |
| 32556 | Protective-serv   | Not-in-family | White | Male   | 0   |
| 32557 | Tech-support      | Wife          | White | Female | 0   |
| 32558 | Machine-op-inspct | Husband       | White | Male   | 0   |
| 32559 | Adm-clerical      | Unmarried     | White | Female | 0   |
| 32560 | Adm-clerical      | Own-child     | White | Male   | 0   |

|       | capital.loss | hours.per.week | native.country | income |
|-------|--------------|----------------|----------------|--------|
| 0     | 4356         | 40             | United-States  | <=50K  |
| 1     | 4356         | 18             | United-States  | <=50K  |
| 2     | 4356         | 40             | United-States  | <=50K  |
| 3     | 3900         | 40             | United-States  | <=50K  |
| 4     | 3900         | 40             | United-States  | <=50K  |
| ...   | ...          | ...            | ...            | ...    |
| 32556 | 0            | 40             | United-States  | <=50K  |
| 32557 | 0            | 38             | United-States  | <=50K  |
| 32558 | 0            | 40             | United-States  | >50K   |
| 32559 | 0            | 40             | United-States  | <=50K  |
| 32560 | 0            | 20             | United-States  | <=50K  |

[32561 rows x 15 columns]

```
[ ]: df.describe
```

```
[ ]: <bound method NDFrame.describe of
education.num    marital.status \
0      90      ?    77053      HS-grad      9      Widowed
1      82  Private  132870      HS-grad      9      Widowed
2      66      ?    186061  Some-college     10      Widowed
3      54  Private  140359      7th-8th      4      Divorced
4      41  Private  264663  Some-college     10      Separated
...  ...
32556  22  Private  310152  Some-college     10      Never-married
32557  27  Private  257302    Assoc-acdm     12  Married-civ-spouse
32558  40  Private  154374      HS-grad      9  Married-civ-spouse
32559  58  Private  151910      HS-grad      9      Widowed
32560  22  Private  201490      HS-grad      9      Never-married
```

|       | occupation        | relationship  | race  | sex    | capital.gain | \ |
|-------|-------------------|---------------|-------|--------|--------------|---|
| 0     | ?                 | Not-in-family | White | Female | 0            |   |
| 1     | Exec-managerial   | Not-in-family | White | Female | 0            |   |
| 2     | ?                 | Unmarried     | Black | Female | 0            |   |
| 3     | Machine-op-inspct | Unmarried     | White | Female | 0            |   |
| 4     | Prof-specialty    | Own-child     | White | Female | 0            |   |
| ...   | ...               | ...           | ...   | ...    | ...          |   |
| 32556 | Protective-serv   | Not-in-family | White | Male   | 0            |   |

|       |                   |           |       |        |   |
|-------|-------------------|-----------|-------|--------|---|
| 32557 | Tech-support      | Wife      | White | Female | 0 |
| 32558 | Machine-op-inspct | Husband   | White | Male   | 0 |
| 32559 | Adm-clerical      | Unmarried | White | Female | 0 |
| 32560 | Adm-clerical      | Own-child | White | Male   | 0 |

|       | capital.loss | hours.per.week | native.country | income |
|-------|--------------|----------------|----------------|--------|
| 0     | 4356         | 40             | United-States  | <=50K  |
| 1     | 4356         | 18             | United-States  | <=50K  |
| 2     | 4356         | 40             | United-States  | <=50K  |
| 3     | 3900         | 40             | United-States  | <=50K  |
| 4     | 3900         | 40             | United-States  | <=50K  |
| ...   | ...          | ...            | ...            | ...    |
| 32556 | 0            | 40             | United-States  | <=50K  |
| 32557 | 0            | 38             | United-States  | <=50K  |
| 32558 | 0            | 40             | United-States  | >50K   |
| 32559 | 0            | 40             | United-States  | <=50K  |
| 32560 | 0            | 20             | United-States  | <=50K  |

[32561 rows x 15 columns]>

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation            32561 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country        32561 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
[ ]: df[df == '?'] = np.nan
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             30725 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation            30718 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country       31978 non-null  object
14  income               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
[ ]: df.isnull().sum()
```

```

[ ]: age                0
workclass             1836
fnlwgt                0
education             0
education.num         0
marital.status        0
occupation            1843
relationship           0
race                  0
sex                   0
capital.gain          0
capital.loss          0
hours.per.week        0
native.country        583
income                0
dtype: int64

```

```

[ ]: max_category = df['workclass'].value_counts().idxmax()
df['workclass'].fillna(max_category, inplace=True)
max_category = df['occupation'].value_counts().idxmax()
df['occupation'].fillna(max_category, inplace=True)
max_category = df['native.country'].value_counts().idxmax()

```




```
df['native.country'].fillna(max_category, inplace=True)
```

```
[ ]: df.isnull().sum()
```

```
[ ]: age                0
      workclass         0
      fnlwgt            0
      education         0
      education.num     0
      marital.status    0
      occupation        0
      relationship      0
      race              0
      sex               0
      capital.gain      0
      capital.loss      0
      hours.per.week    0
      native.country    0
      income            0
      dtype: int64
```

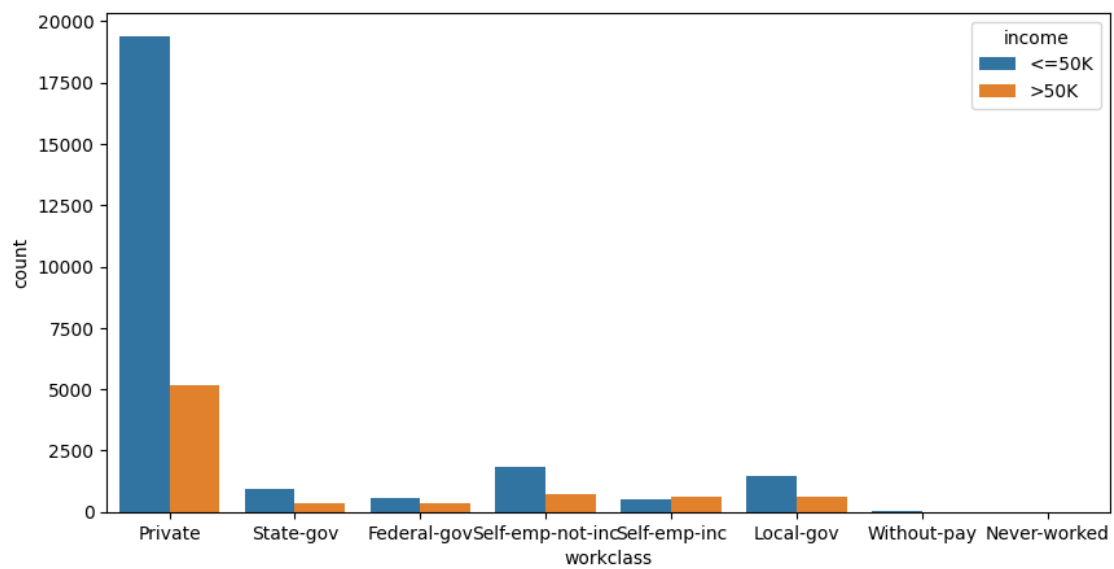
```
[ ]: sb.histplot(df, x='age', hue='income', bins= 32)
```

```
[ ]: <Axes: xlabel='age', ylabel='Count'>
```




```
[ ]: fig=plt.figure(figsize=(10,5))
sb.countplot(data = df, x = 'workclass', hue = 'income')
```

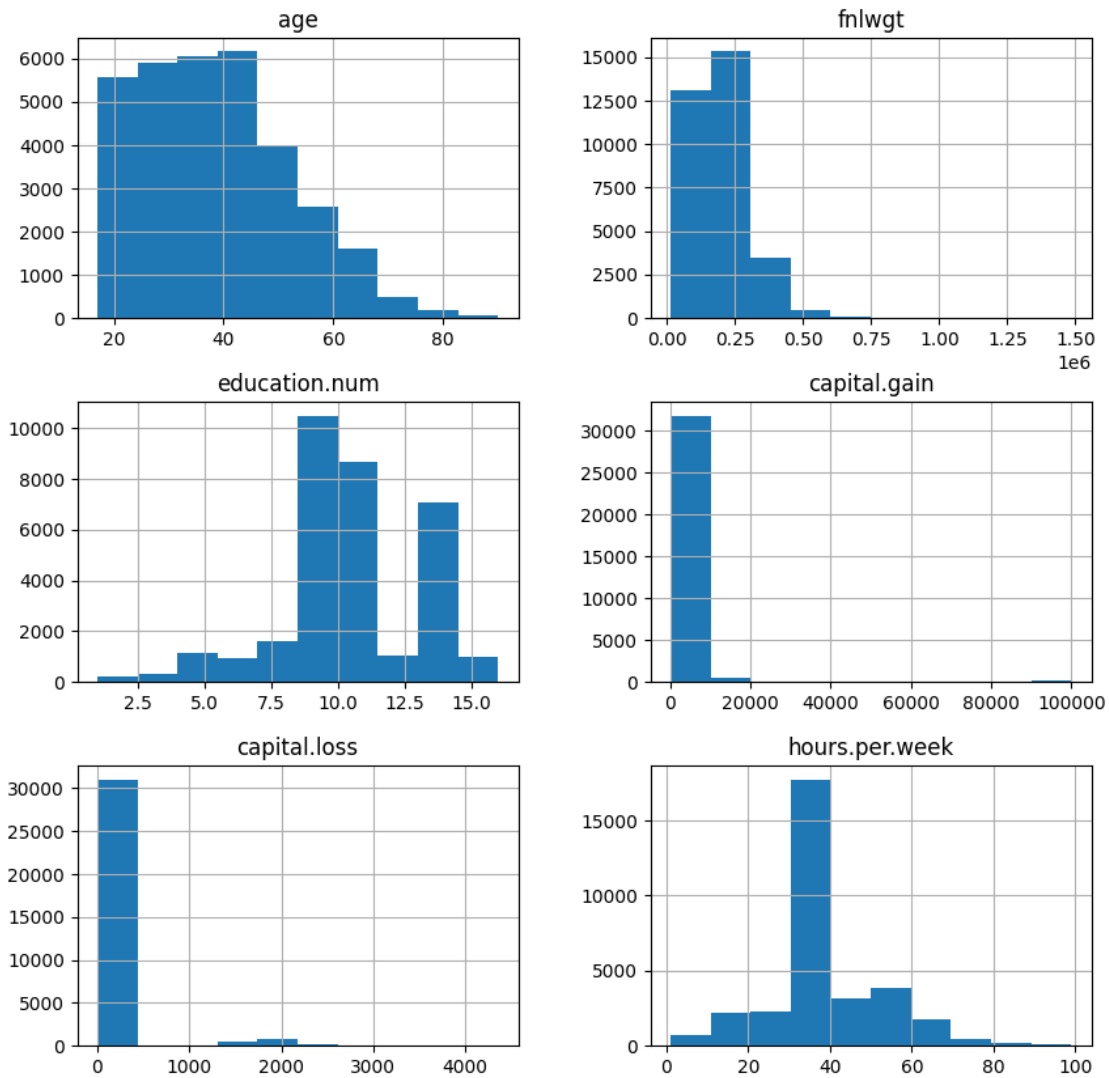
```
[ ]: <Axes: xlabel='workclass', ylabel='count'>
```




```
[ ]: sb.countplot(data = df, x = 'education', hue = 'income')
plt.tick_params(axis='x', rotation=90)
```



```
[ ]: df.hist(bins=10, figsize=(10, 10))
plt.show()
```

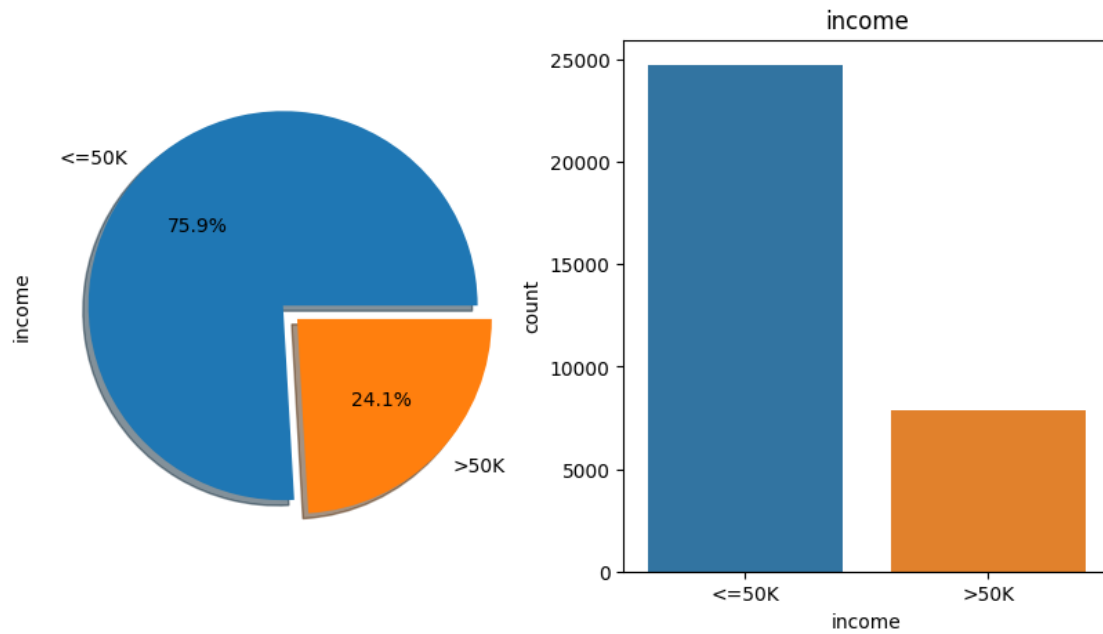


```
[ ]: sb.countplot(x = "race", data=df);
```



```
[ ]: f, ax = plt.subplots(1, 2, figsize=(10, 5))
df['income'].value_counts().plot.pie(explode=[0, 0.1], autopct='%1.1f%%',
    ↪ax=ax[0], shadow=True)
sb.countplot(x='income', data=df, ax=ax[1])
ax[1].set_title('income')
```

```
[ ]: Text(0.5, 1.0, 'income')
```




```
[ ]: corr=df.corr()  
sb.heatmap(corr,annot=True)
```

<ipython-input-279-5b574c6aa484>:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
corr=df.corr()
```

```
[ ]: <Axes: >
```



```
[ ]: df=df.drop(columns='race')
df=df.drop(columns='fnlwgt')
df=df.drop(columns='education.num')
df=df.drop(columns='relationship')
df=df.drop(columns='native.country')
# df=df.drop(columns='marital.status')
```

```
[ ]: from sklearn.preprocessing import LabelEncoder
```

```
[ ]: labelencoder_x=LabelEncoder()
df["workclass"]=labelencoder_x.fit_transform(df["workclass"])
df["education"]=labelencoder_x.fit_transform(df["education"])
# df["relationship"]=labelencoder_x.fit_transform(df["relationship"])
df["occupation"]=labelencoder_x.fit_transform(df["occupation"])
df["sex"]=labelencoder_x.fit_transform(df["sex"])
# df["native.country"]=labelencoder_x.fit_transform(df["native.country"])
df["income"]=labelencoder_x.fit_transform(df["income"])
df["marital.status"]=labelencoder_x.fit_transform(df["marital.status"])
```

```
[ ]: x=df.drop("income",axis=1)
      y=df["income"]
```

```
[ ]: df.head(20)
```

```
[ ]:
      age  workclass  education  marital.status  occupation  sex  capital.gain  \
0      90         3         11             6           9      0           0
1      82         3         11             6           3      0           0
2      66         3         15             6           9      0           0
3      54         3          5             0           6      0           0
4      41         3         15             5           9      0           0
5      34         3         11             0           7      0           0
6      38         3          0             5           0      1           0
7      74         6         10             4           9      0           0
8      68         0         11             0           9      0           0
9      41         3         15             4           2      1           0
10     45         3         10             0           9      0           0
11     38         5         14             4           9      1           0
12     52         3          9             6           7      0           0
13     32         3         12             5           3      1           0
14     51         3         10             4           9      1           0
15     46         3         14             0           9      1           0
16     45         3          1             0          13      1           0
17     57         3         12             0           3      1           0
18     22         3          7             4           5      1           0
19     34         3          9             5          11      1           0
```

```
      capital.loss  hours.per.week  income
0           4356           40         0
1           4356           18         0
2           4356           40         0
3           3900           40         0
4           3900           40         0
5           3770           45         0
6           3770           40         0
7           3683           20         1
8           3683           40         0
9           3004           60         1
10          3004           35         1
11          2824           45         1
12          2824           20         1
13          2824           55         1
14          2824           40         1
15          2824           40         1
16          2824           76         1
17          2824           50         1
18          2824           40         1
```



```
[ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=0,test_size=0.3)
```

```
[ ]: from sklearn import tree
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report
      from sklearn.ensemble import RandomForestClassifier
```

```
[ ]: dt=DecisionTreeClassifier(max_depth=5)
      dt.fit(x_train,y_train)
```

```
[ ]: DecisionTreeClassifier(max_depth=5)
```

```
[ ]: features = list(df.columns[1:])
      features
```

```
[ ]: ['workclass',
      'education',
      'marital.status',
      'occupation',
      'sex',
      'capital.gain',
      'capital.loss',
      'hours.per.week',
      'income']
```

```
[ ]: y_dtp=dt.predict(x_test)
```

```
[ ]: print(classification_report(y_test,y_dtp))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.81      | 1.00   | 0.90     | 7410    |
| 1            | 0.98      | 0.27   | 0.43     | 2359    |
| accuracy     |           |        | 0.82     | 9769    |
| macro avg    | 0.90      | 0.64   | 0.66     | 9769    |
| weighted avg | 0.85      | 0.82   | 0.78     | 9769    |

```
[ ]: rf=RandomForestClassifier(random_state=1)
      rf.fit(x_train,y_train)
      y_rfp=rf.predict(x_test)
```

```
[ ]: print('Random Forest : ',accuracy_score(y_test,y_rfp)*100)
```

Random Forest : 84.58388780837342

```
[ ]: !pip install my-package
```

Requirement already satisfied: my-package in /usr/local/lib/python3.10/dist-packages (0.0.0)

```
[ ]: !pip install pydotplus
```

Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-packages (2.0.2)

Requirement already satisfied: pyparsing>=2.0.1 in /usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)

```
[ ]: from IPython.display import Image
from six import StringIO
from sklearn.tree import export_graphviz
import pydotplus,graphviz
```

```
[ ]: !pip install graphviz
```

Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (0.20.1)

```
[ ]: dot_data = StringIO()
export_graphviz(dt, out_file=dot_data,
feature_names=features, filled=True,rounded=True)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

```
[ ]:
```

