

Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:

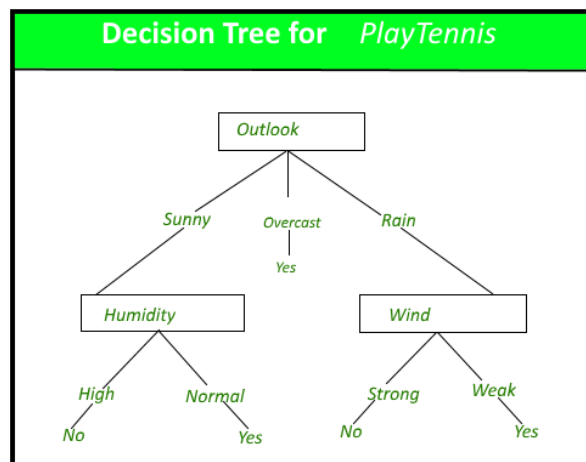


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:



Vidyavardhini's College of Engineering & Technology
Department of Computer Engineering

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras,



Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Conclusion:

1. Discuss about the how categorical attributes have been dealt with during data pre-processing.

Categorical attributes, also known as categorical variables or features, are those that represent discrete and distinct categories or groups. These attributes are common in data across various domains, such as gender, color, city names, and more. Label encoding is a simple technique where each category is assigned a unique integer label. For example, if you have a "color" attribute with categories like "red," "blue," and "green," you can encode them as 0, 1, and 2, respectively. Here Label encoder is used to convert categorical to numerical.

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.

Hyperparameter tuning is a critical step in optimizing the performance of a decision tree model. Decision trees have several hyperparameters that control their structure and behavior, and tuning these hyperparameters can help improve the model's accuracy, generalization, and robustness. Max depth controls the maximum depth of the tree. A deeper tree can capture more complex relationships in the data but is prone to overfitting. We have used max depth as 5.

3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

Accuracy obtained in the decision tree model is 85%. A confusion matrix is a tabular representation used in machine learning to evaluate the performance of a classification model, especially for binary classification problems. Precision Obtain is 0.98 ,recall obtained is 0.27 and f1 score is 0.43.

adultcensus

Adult Census Income Dataset

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
```

```
[ ]: df=pd.read_csv("/content/adult.csv",error_bad_lines=False, engine="python")
```

<ipython-input-265-8dbf926b9719>:1: FutureWarning: The error_bad_lines argument has been deprecated and will be removed in a future version. Use on_bad_lines in the future.

```
df=pd.read_csv("/content/adult.csv",error_bad_lines=False, engine="python")
```

```
[ ]: print(df)
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
0	90	?	77053	HS-grad	9	Widowed	
1	82	Private	132870	HS-grad	9	Widowed	
2	66	?	186061	Some-college	10	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
...	
32556	22	Private	310152	Some-college	10	Never-married	
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	
32558	40	Private	154374	HS-grad	9	Married-civ-spouse	
32559	58	Private	151910	HS-grad	9	Widowed	
32560	22	Private	201490	HS-grad	9	Never-married	

	occupation	relationship	race	sex	capital.gain	\
0		?	Not-in-family	White	Female	0
1	Exec-managerial	Not-in-family	White	Female		0
2		?	Unmarried	Black	Female	0
3	Machine-op-inspct	Unmarried	White	Female		0

4	Prof-specialty	Own-child	White	Female	0
...
32556	Protective-serv	Not-in-family	White	Male	0
32557	Tech-support	Wife	White	Female	0
32558	Machine-op-inspct	Husband	White	Male	0
32559	Adm-clerical	Unmarried	White	Female	0
32560	Adm-clerical	Own-child	White	Male	0

	capital.loss	hours.per.week	native.country	income
0	4356	40	United-States	<=50K
1	4356	18	United-States	<=50K
2	4356	40	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K
...
32556	0	40	United-States	<=50K
32557	0	38	United-States	<=50K
32558	0	40	United-States	>50K
32559	0	40	United-States	<=50K
32560	0	20	United-States	<=50K

[32561 rows x 15 columns]

```
[ ]: df.describe
```

```
[ ]: <bound method NDFrame.describe of
education.num    marital.status \
0      90      ?    77053      HS-grad      9      Widowed
1      82  Private  132870      HS-grad      9      Widowed
2      66      ?    186061  Some-college     10      Widowed
3      54  Private  140359      7th-8th      4      Divorced
4      41  Private  264663  Some-college     10      Separated
...
32556    22  Private  310152  Some-college     10      Never-married
32557    27  Private  257302    Assoc-acdm     12  Married-civ-spouse
32558    40  Private  154374      HS-grad      9  Married-civ-spouse
32559    58  Private  151910      HS-grad      9      Widowed
32560    22  Private  201490      HS-grad      9      Never-married
```

	occupation	relationship	race	sex	capital.gain	\
0	?	Not-in-family	White	Female	0	
1	Exec-managerial	Not-in-family	White	Female	0	
2	?	Unmarried	Black	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
...	
32556	Protective-serv	Not-in-family	White	Male	0	

32557	Tech-support	Wife	White	Female	0
32558	Machine-op-inspct	Husband	White	Male	0
32559	Adm-clerical	Unmarried	White	Female	0
32560	Adm-clerical	Own-child	White	Male	0

	capital.loss	hours.per.week	native.country	income
0	4356	40	United-States	<=50K
1	4356	18	United-States	<=50K
2	4356	40	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K
...
32556	0	40	United-States	<=50K
32557	0	38	United-States	<=50K
32558	0	40	United-States	>50K
32559	0	40	United-States	<=50K
32560	0	20	United-States	<=50K

[32561 rows x 15 columns]>

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
5   marital.status         32561 non-null  object
6   occupation             32561 non-null  object
7   relationship           32561 non-null  object
8   race                   32561 non-null  object
9   sex                    32561 non-null  object
10  capital.gain           32561 non-null  int64
11  capital.loss           32561 non-null  int64
12  hours.per.week         32561 non-null  int64
13  native.country         32561 non-null  object
14  income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
[ ]: df[df == '?'] = np.nan
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             30725 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation            30718 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week       32561 non-null  int64
13  native.country       31978 non-null  object
14  income               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
[ ]: df.isnull().sum()
```

```

[ ]: age                0
workclass             1836
fnlwgt                0
education             0
education.num         0
marital.status        0
occupation            1843
relationship           0
race                  0
sex                   0
capital.gain           0
capital.loss           0
hours.per.week        0
native.country        583
income                 0
dtype: int64

```

```

[ ]: max_category = df['workclass'].value_counts().idxmax()
df['workclass'].fillna(max_category, inplace=True)
max_category = df['occupation'].value_counts().idxmax()
df['occupation'].fillna(max_category, inplace=True)
max_category = df['native.country'].value_counts().idxmax()

```




```
df['native.country'].fillna(max_category, inplace=True)
```

```
[ ]: df.isnull().sum()
```

```
[ ]: age                0
      workclass         0
      fnlwgt            0
      education         0
      education.num     0
      marital.status    0
      occupation        0
      relationship      0
      race              0
      sex               0
      capital.gain      0
      capital.loss      0
      hours.per.week    0
      native.country    0
      income            0
      dtype: int64
```


```
[ ]: sb.histplot(df, x='age', hue='income', bins= 32)
```

```
[ ]: <Axes: xlabel='age', ylabel='Count'>
```




```
[ ]: fig=plt.figure(figsize=(10,5))
sb.countplot(data = df, x = 'workclass', hue = 'income')
```

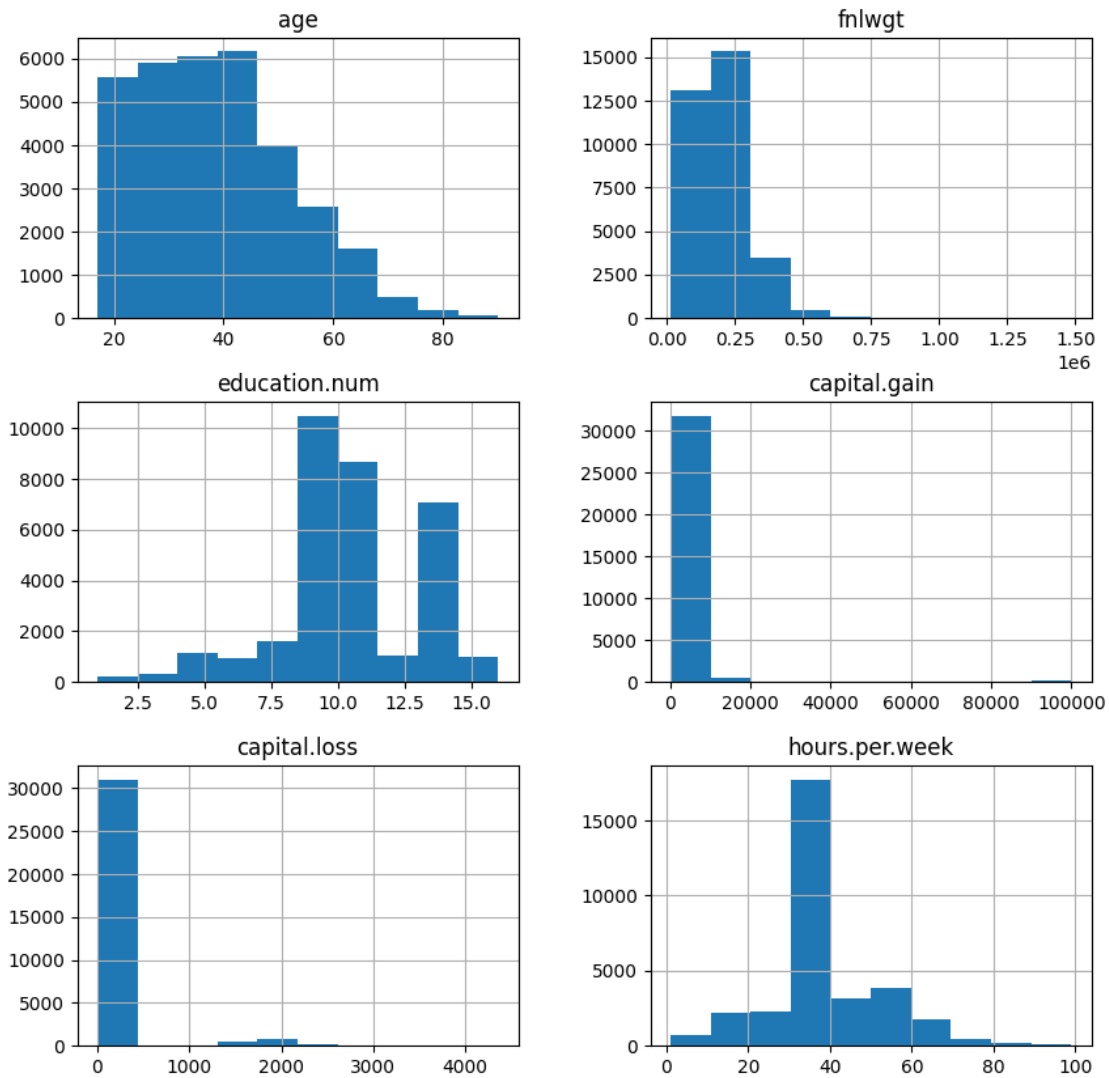
```
[ ]: <Axes: xlabel='workclass', ylabel='count'>
```




```
[ ]: sb.countplot(data = df, x = 'education', hue = 'income')
plt.tick_params(axis='x', rotation=90)
```



```
[ ]: df.hist(bins=10, figsize=(10, 10))
plt.show()
```

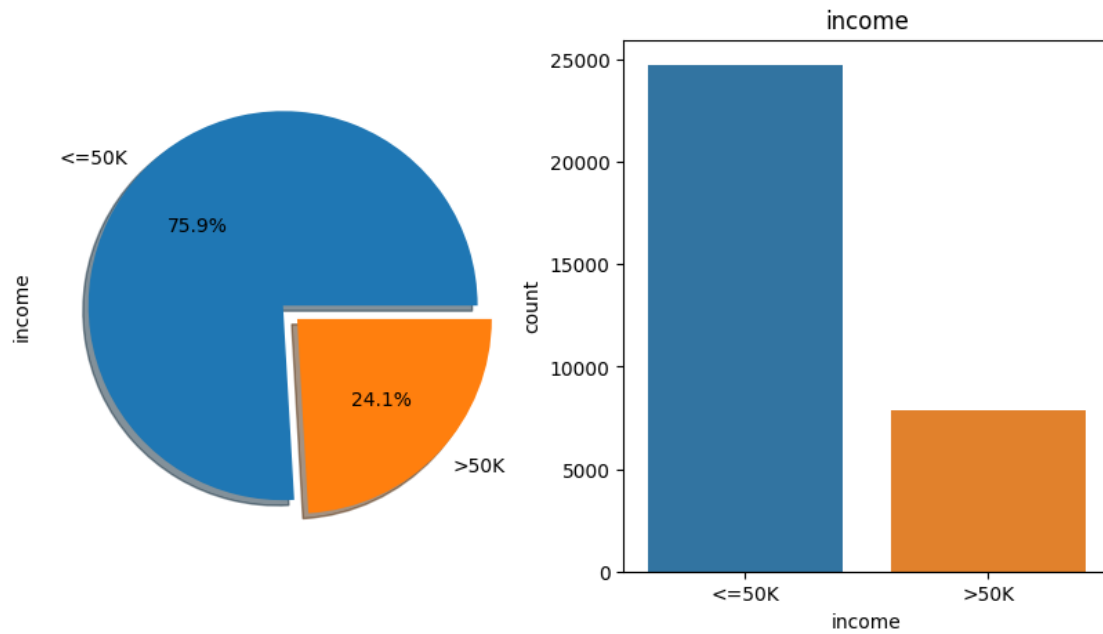


```
[ ]: sb.countplot(x = "race", data=df);
```



```
[ ]: f, ax = plt.subplots(1, 2, figsize=(10, 5))
df['income'].value_counts().plot.pie(explode=[0, 0.1], autopct='%1.1f%%',
    ↪ax=ax[0], shadow=True)
sb.countplot(x='income', data=df, ax=ax[1])
ax[1].set_title('income')
```

```
[ ]: Text(0.5, 1.0, 'income')
```




```
[ ]: corr=df.corr()  
sb.heatmap(corr,annot=True)
```

<ipython-input-279-5b574c6aa484>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr=df.corr()
```

```
[ ]: <Axes: >
```



```
[ ]: df=df.drop(columns='race')
df=df.drop(columns='fnlwgt')
df=df.drop(columns='education.num')
df=df.drop(columns='relationship')
df=df.drop(columns='native.country')
# df=df.drop(columns='marital.status')
```

```
[ ]: from sklearn.preprocessing import LabelEncoder
```

```
[ ]: labelencoder_x=LabelEncoder()
df["workclass"]=labelencoder_x.fit_transform(df["workclass"])
df["education"]=labelencoder_x.fit_transform(df["education"])
# df["relationship"]=labelencoder_x.fit_transform(df["relationship"])
df["occupation"]=labelencoder_x.fit_transform(df["occupation"])
df["sex"]=labelencoder_x.fit_transform(df["sex"])
# df["native.country"]=labelencoder_x.fit_transform(df["native.country"])
df["income"]=labelencoder_x.fit_transform(df["income"])
df["marital.status"]=labelencoder_x.fit_transform(df["marital.status"])
```

```
[ ]: x=df.drop("income",axis=1)
      y=df["income"]
```

```
[ ]: df.head(20)
```

```
[ ]:
      age  workclass  education  marital.status  occupation  sex  capital.gain  \
0      90         3         11             6           9      0           0
1      82         3         11             6           3      0           0
2      66         3         15             6           9      0           0
3      54         3          5             0           6      0           0
4      41         3         15             5           9      0           0
5      34         3         11             0           7      0           0
6      38         3          0             5           0      1           0
7      74         6         10             4           9      0           0
8      68         0         11             0           9      0           0
9      41         3         15             4           2      1           0
10     45         3         10             0           9      0           0
11     38         5         14             4           9      1           0
12     52         3          9             6           7      0           0
13     32         3         12             5           3      1           0
14     51         3         10             4           9      1           0
15     46         3         14             0           9      1           0
16     45         3          1             0          13      1           0
17     57         3         12             0           3      1           0
18     22         3          7             4           5      1           0
19     34         3          9             5          11      1           0
```

```
      capital.loss  hours.per.week  income
0           4356           40         0
1           4356           18         0
2           4356           40         0
3           3900           40         0
4           3900           40         0
5           3770           45         0
6           3770           40         0
7           3683           20         1
8           3683           40         0
9           3004           60         1
10          3004           35         1
11          2824           45         1
12          2824           20         1
13          2824           55         1
14          2824           40         1
15          2824           40         1
16          2824           76         1
17          2824           50         1
18          2824           40         1
```



```
[ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=0,test_size=0.3)
```

```
[ ]: from sklearn import tree
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report
      from sklearn.ensemble import RandomForestClassifier
```

```
[ ]: dt=DecisionTreeClassifier(max_depth=5)
      dt.fit(x_train,y_train)
```

```
[ ]: DecisionTreeClassifier(max_depth=5)
```

```
[ ]: features = list(df.columns[1:])
      features
```

```
[ ]: ['workclass',
      'education',
      'marital.status',
      'occupation',
      'sex',
      'capital.gain',
      'capital.loss',
      'hours.per.week',
      'income']
```

```
[ ]: y_dtp=dt.predict(x_test)
```

```
[ ]: print(classification_report(y_test,y_dtp))
```

	precision	recall	f1-score	support
0	0.81	1.00	0.90	7410
1	0.98	0.27	0.43	2359
accuracy			0.82	9769
macro avg	0.90	0.64	0.66	9769
weighted avg	0.85	0.82	0.78	9769

```
[ ]: rf=RandomForestClassifier(random_state=1)
      rf.fit(x_train,y_train)
      y_rfp=rf.predict(x_test)
```

```
[ ]: print('Random Forest : ',accuracy_score(y_test,y_rfp)*100)
```

Random Forest : 84.58388780837342

```
[ ]: !pip install my-package
```

```
Requirement already satisfied: my-package in /usr/local/lib/python3.10/dist-packages (0.0.0)
```

```
[ ]: !pip install pydotplus
```

```
Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-packages (2.0.2)
```

Requirement already satisfied: pyparsing>=2.0.1 in
/usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)

```
[ ]: from IPython.display import Image
      from six import StringIO
      from sklearn.tree import export_graphviz
      import pydotplus, graphviz
```

```
[ ]: !pip install graphviz
```

```
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (0.20.1)
```

```
[ ]: dot_data = StringIO()
      export_graphviz(dt, out_file=dot_data,
                      feature_names=features, filled=True, rounded=True)
      graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
      Image(graph.create_png())
```

[]:

